

ADM Project

Maurizio Meschi & Nicolò Trebino

January 30, 2025

Email: maurizio.meschi@gmail.com, nicolo.trebino@gmail.com

GitHub: [@Maurizio-Meschi](#), [@nicolotrebino](#)

Repo: [GitHub Repository](#)

Contents

1	Introduction	2
2	System requirements	2
3	ER diagram	3
4	Workload	3
5	Aggregates methodology	4
5.1	Write the query in a formal and non-ambiguous way	4
5.2	Annotated the ER schema	5
5.3	Aggregates	5
6	Cassandra	6
6.1	Schema design	6
6.1.1	Book aggregate	6
6.1.2	Author aggregate	7
6.1.3	User aggregate	8
6.2	Query Modeling in CQL	8
7	Neo4j	10
7.1	Schema design	10
7.2	Query modeling in Cypher	11
8	Why Cassandra?	12
9	Cassandra configuration	12
10	Cassandra logical schema	13
10.1	Keyspace	13
10.2	Types	13
10.3	Tables	13
10.4	Copy the dataset from CSV files	13
11	Cassandra instance	14
12	Workload implementation	15
12.1	Q1	15
12.2	Q2	15
12.3	Q3	15
12.4	Q4	16
12.5	Q5	16
12.6	Q6	17
12.7	Q7	17
12.8	Q8	18
13	Conclusions	18

1 Introduction

As passionate people who read books, we have chosen to take inspiration from this [dataset](#) from [kaggle.com](#).

However, since it does not fit our application domain quite well, we have used [Hardcover API](#) to retrieve all the needed data and create the data set.

Our domain application consists of an on-line catalog where USERS can REVIEW and rate all the books and share their thoughts with a huge community of readers around the world.

Our application aims to help people find the perfect book that matches their preferences, and we chose a workload that includes the main operations that a normal user usually does in our catalog.

A possible future implementation can be a recommendation system that is very important in all modern social network applications that, based on what a user reads or likes, can recommend books or users to our customers, strengthening the interconnections between users, improving the overall impact and effectiveness of our application.

2 System requirements

Our application is suitable for a transactional, read/write intensive scenario since users can search for specific books and information and write reviews and scores.

Batch processing is not the focus of our application, even if our catalog automatically performs the average of the ratings of a particular book when a user lends on the book page: an action that requires batch access in the database.

We want our system to be scalable to provide for an ever-increasing number of requests in the future, expecting an ever-growing number of users. We also need high availability and high throughput to provide all requests, both read and write or update, that our users may make.

Partition tolerance is also a key point in our system since we need the system to continue to work even after a network partition or failure, so we have to rely, following the CAP theorem, on an eventual consistency level, to guarantee that, given enough time and no new updates, all replicas of data will eventually converge to the same state.

3 ER diagram

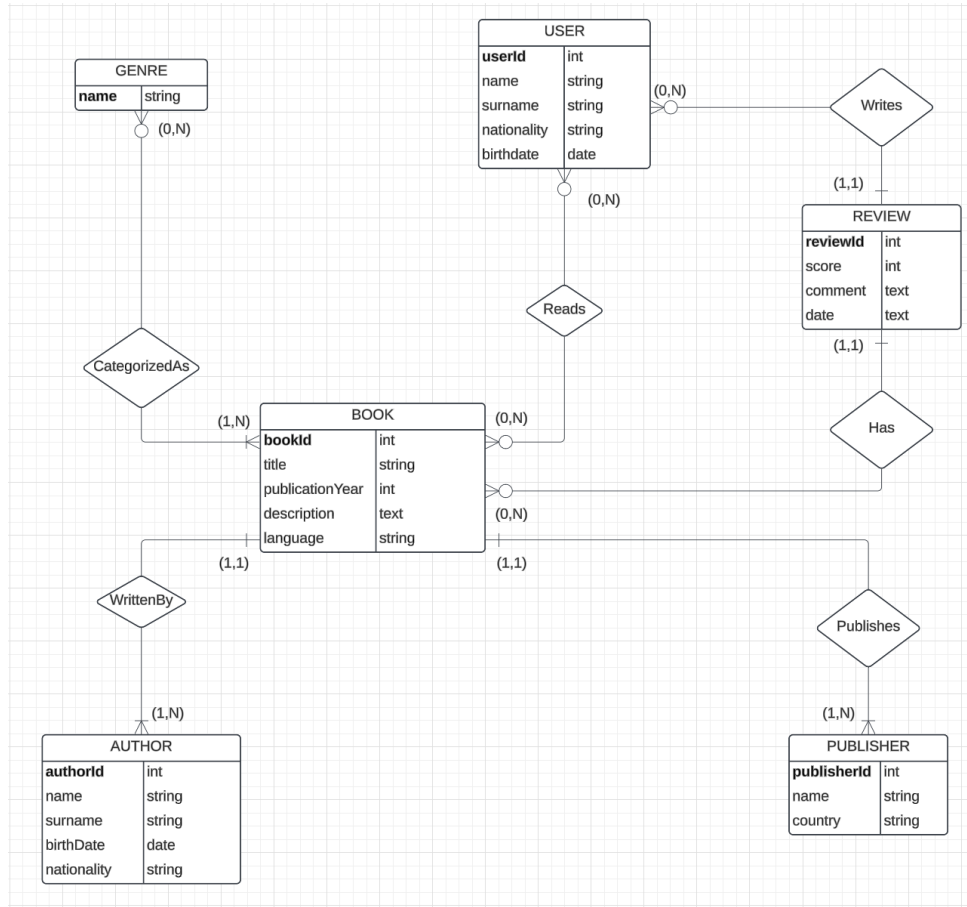


Figure 1: Library relational model

4 Workload

- Q1:** Given the book with ID 218619, determine the average score and all the book information.
- Q2:** Determine all books written by 737939 of the genre “Science fiction”.
- Q3:** Find all books written by the author with ID 206880 along with their publication year.
- Q4:** Find all books published in 2017 in English.
- Q5:** Provide all the review information for the book with ID 350392, including the users who wrote them.
- Q6:** Given the user with ID 54, determine all the information about the books read by them and their genres.
- Q7:** Provide all the books published in the last year along with their publishers.
- Q8:** Show 10 books for the home page.

5 Aggregates methodology

5.1 Write the query in a formal and non-ambiguous way

Q1. (BOOK, [BOOK(bookId)_!], [BOOK_!, REVIEW(score)_H])

E = BOOK

LS = [BOOK(bookId)_!]

LP = [BOOK_!, REVIEW(score)_H]

Q2. (AUTHOR, [GENRE_WbCa, AUTHOR(authorId)_!], [BOOK(bookId)_Wb])

E = AUTHOR

LS = [GENRE_WbCa, AUTHOR(authorId)_!]

LP = [BOOK(bookId)_Wb]

Q3. (AUTHOR, [AUTHOR(authorId)_!], [BOOK(bookId, publicationYear)_Wb])

E = AUTHOR

LS = [AUTHOR(authorId)_!]

LP = [BOOK(bookId, publicationYear)_Wb]

Q4. (BOOK, [BOOK(publicationYear, language)_!], [BOOK(bookId)_!])

E = BOOK

LS = [BOOK(publicationYear, language)_!]

LP = [BOOK(bookId)_!]

Q5. (BOOK, [BOOK(bookId)_!], [REVIEW_H, USER(userId)_HW])

E = BOOK

LS = [BOOK(bookId)_!]

LP = [REVIEW_H, USER(userId)_HW]

Q6. (USER, [USER(userId)_!], [BOOK_R, GENRE_RCa])

E = USER

LS = [USER(userId)_!]

LP = [BOOK_R, GENRE_RCa]

Q7. (BOOK, [BOOK(publicationYear)_!], [BOOK(bookId)_!, PUBLISHER(publisherId)_P])

E = BOOK

LS = [BOOK(publicationYear)_!]

LP = [BOOK(bookId)_!, PUBLISHER(publisherId)_P]

Q8. (BOOK, [], [BOOK(bookId)_!])

E = BOOK

LS = []

LP = [BOOK(bookId)_!]

5.2 Annotated the ER schema

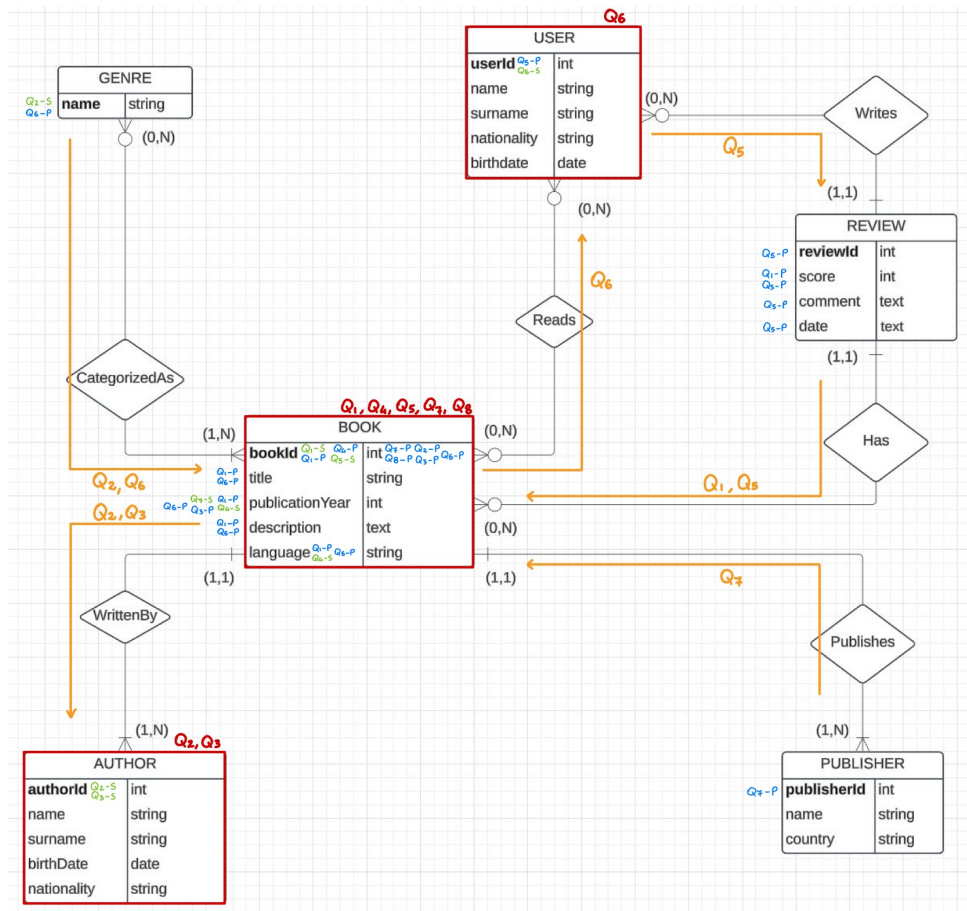


Figure 2: Annotated entity relationship diagram

5.3 Aggregates

```
Book: {
  bookId, title, publicationYear, description, language,
  reviews: [{ reviewId, score, comment, date, userId }],
  publisherId
}
```

```
Author: {
  authorId,
  books: [{
    bookId, publicationYear,
    genres: [{ name }]
  }]
}
```

```
User: {
  userId,
  books: [{
    bookId, title, publicationYear, description, language,
    genres: [{ name }]
  }]
}
```

6 Cassandra

6.1 Schema design

6.1.1 Book aggregate

Q1. (BOOK, [BOOK(bookId)_!], [BOOK_!, REVIEW(score)_H])

Q4. (BOOK, [BOOK(publicationYear, language)_!], [BOOK(bookId)_!])

Q5. (BOOK, [BOOK(bookId)_!], [REVIEW_H, USER(userId)_HW])

Q7. (BOOK, [BOOK(publicationYear)_!],[BOOK(bookId)_!, PUBLISHER(publisherId)_P])

Q8. (BOOK, [], [BOOK(bookId)_!])

Since the **LS** sets are disjoint, we define the following tables:

- Q1, Q5, Q8: Table with BOOK(bookId) as unique **partition key** and **primary key**.
- Q4, Q7: Table with BOOK(publicationYear) as **partition key** and BOOK(language) as **clustering column** together with BOOK(bookId) (identifier of the aggregate).
- Define a new type for the nested attribute **reviews** inside the aggregate.

Defining a new type for reviews:

```
CREATE TYPE IF NOT EXISTS Review_t (  
    reviewId    int,  
    score       int,  
    comment     text,  
    date        varchar,  
    userId      int  
);
```

Tables for BOOK:

```
CREATE TABLE IF NOT EXISTS Book1 (  
    bookId      int PRIMARY KEY,  
    title       varchar,  
    publicationYear int,  
    description  text,  
    language    varchar,  
    reviews     set<frozen<Review_t>>,  
    publisherId int  
);
```

```
CREATE TABLE IF NOT EXISTS Book2 (  
    bookId      int,  
    title       varchar,  
    publicationYear int,  
    description  text,  
    language    varchar,  
    reviews     set<frozen<Review_t>>,  
    publisherId int,  
    PRIMARY KEY (publicationYear, language, bookId)  
);
```

6.1.2 Author aggregate

Q2. (AUTHOR, [GENRE_WbCa, AUTHOR(authorId)_!], [BOOK(bookId)_Wb])

Q3. (AUTHOR, [AUTHOR(authorId)_!], [BOOK(bookId, publicationYear)_Wb])

For the Author aggregate, given these two queries from the workload, we must pay attention.

There are no problems for Q3 since we can just create a table with AuthorId as the partition key (primary key), instead for Q2, due to the logical architecture of Cassandra we cannot perform the query.

In fact, for Q2 we should have to create a table with GENRE(name) as a clustering column, but this is not possible because GENRE(name) is a nested attribute in the aggregate, and that means we cannot specify conditions on GENRE(name), neither creating some index nor specifying the ALLOWING FILTERING clause!

In this case, we should modify the aggregation in such a way we can perform the query efficiently but if we chose GENRE as the new aggregation entity, the same problem would persist on AUTHOR(AuthorId) because the cardinality is always **many to many**, so the problem of nested attribute persists.

Hence, we create a single table with AUTHOR(AuthorId) as the **partition key (primary key)** and to perform Q2 we must retrieve, at the application level, all the information that we need. In particular, we can obtain all the books written by a specific author, but we can filter by the genre only at the application level.

In short:

- for **Q3**, we can use AUTHOR(authorId) as **partition key**.
- **Q2** cannot directly be implemented in Cassandra due to the nested attribute GENRE(name). We solve this by handling filtering at the **application level**.
- So AUTHOR(authorId) will be both **partition key** and **primary key**.

Defining a new type for books:

```
CREATE TYPE IF NOT EXISTS Book_at (  
    bookId          int,  
    publicationYear int,  
    genres          set<text>  
);
```

Table for AUTHOR:

```
CREATE TABLE IF NOT EXISTS Author (  
    authorId      int PRIMARY KEY,  
    books         set<frozen<Book_at>>  
);
```


6.1.3 User aggregate

Q6. (USER, [USER(userId)_!], [BOOK_R, GENRE_RCa])

Since we have only one query, we create a single table with USER(userId) as the **partition key** (primary key) to perform this query efficiently.

Defining a new type for books:

```
CREATE TYPE IF NOT EXISTS Book_ut (  
    bookId          int,  
    title           varchar,  
    publicationYear int,  
    description     text,  
    language        varchar,  
    genres          set<text>  
);
```

Table for USER:

```
CREATE TABLE IF NOT EXISTS User (  
    userId    int PRIMARY KEY,  
    books     set<frozen<Book_ut>>  
);
```

6.2 Query Modeling in CQL

Q1. Given the book with ID 218619, determine the average score and all the book information.

```
SELECT *  
FROM   Book1  
WHERE  bookId = 218619
```

The score is a nested attribute inside **reviews**, so we can only get the reviews of the book and compute the average of the score at the application level.

Q2. Determine all books written by 737939 of the genre “Science fiction”.

```
SELECT books  
FROM   Author  
WHERE  authorId = 737939
```

As we already said, we cannot filter by genre, since it is a nested attribute. We have to perform this query to retrieve all the books written by the author 737939 and then filter by the genre at the application level.

Q3. Find all books written by the author with ID 206880 along with their publication year.

```
SELECT books
FROM Author
WHERE authorId = 206880
```

In this case, we also get the genres of all the books, because **books** is a nested attribute. If we are not interested in genres, we have to project just what we need at the application level.

Q4. Find all books published in 2017 in English.

```
SELECT bookId
FROM Book2
WHERE publicationYear = 2017 AND language = 'English'
```

Q5. Provide all the review information for the book with ID 350392, including the users who wrote them.

```
SELECT reviews
FROM Book1
WHERE bookId = 350392
```

Q6. Given the user ID 54, determine all the information about the books read by them and their genres.

```
SELECT books
FROM User
WHERE userId = 54
```

Q7. Provide all the books published in the last year along with their publishers.

```
SELECT bookId, publisherId
FROM Book2
WHERE publicationYear = 2024
```

Q8. Show 10 books for the home page.

```
SELECT bookId
FROM Book1
LIMIT 10
```

7 Neo4j

7.1 Schema design

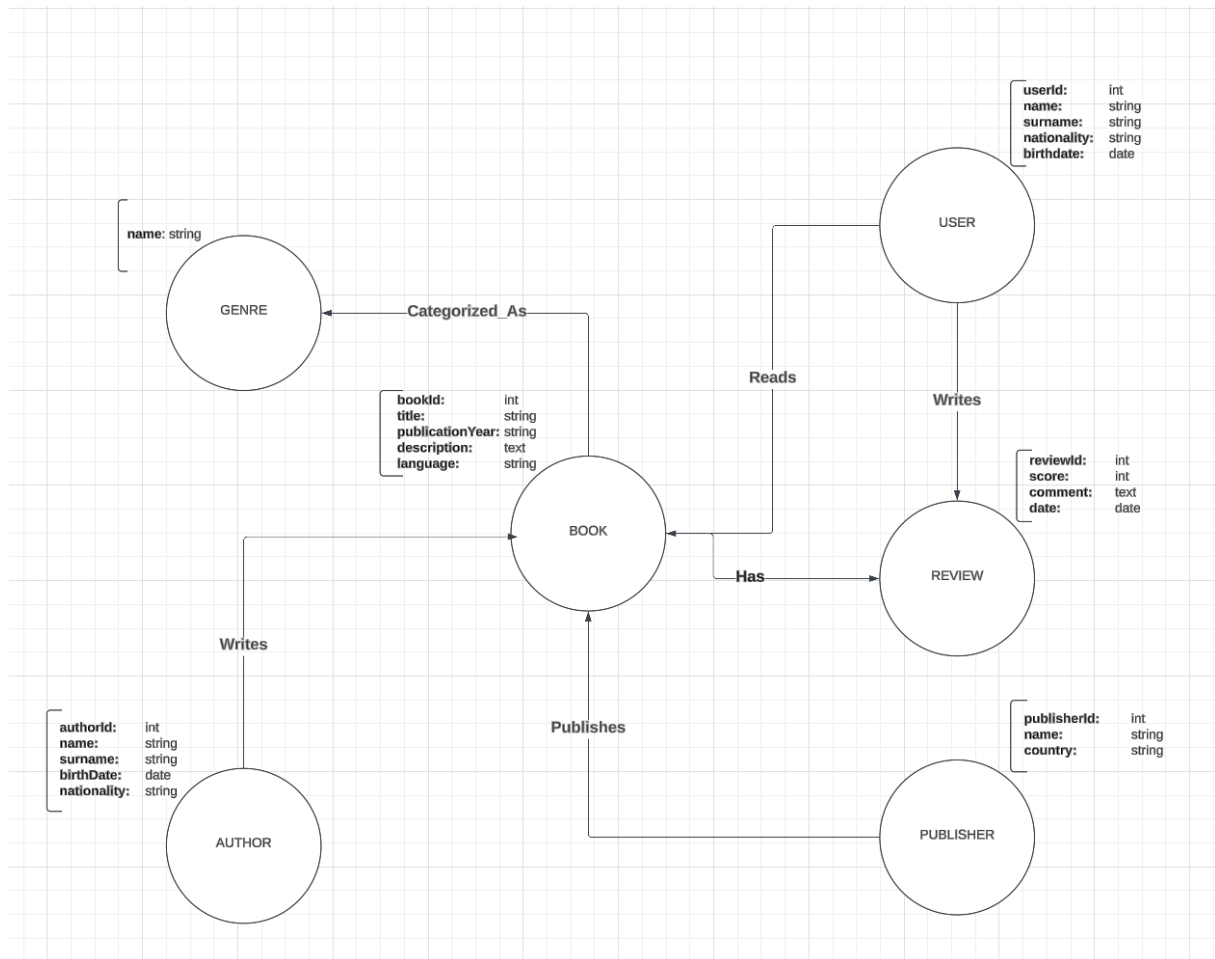


Figure 3: Neo4j graph schema

Graph-based NoSQL systems, such as Neo4j, are designed to handle complex, highly interconnected data.

Starting from our ER diagram we chose to model the graph in such a way we can handle all the queries in an efficient way.

There are many different ways to design the graph. For example, one might consider merging **genre** into the book node as an attribute. However, this would not be a good idea for our application, as each book can belong to multiple genres. Keeping **genre** as a separate node provides greater flexibility, allows for more complex queries, and prevents data duplication since the same genre node can be linked to multiple books at the same time.

One might also think to insert years or dates as labels for the relationship, but for our workload, we chose to construct the model avoiding useless traversal, even if they are very efficient in graph DB.

7.2 Query modeling in Cypher

Q1. Given the book with ID 1598407, determine the average score and all the book information.

```
MATCH (b:Book{bookId:1598407})-[:Has]->(r:Review)
RETURN b, AVG(r.score)
```

Q2. Determine all the books written by 110620 of the genre “Fiction”.

```
MATCH (:Genre{name:"Fiction"})<-[:Categorized_As]-(b:Book)
      <-[:Writes]-(a:Author{authorId:110620})
RETURN b.bookId
```

Q3. Find all the books written by the author with ID 402411, along with their publication year.

```
MATCH (b: Book)<-[:Writes]-(a:Author{authorId: 402411})
RETURN b.bookId, b.publicationYear
```

Q4. Find all the books published in 2008 in English.

```
MATCH (b:Book)
WHERE b.publicationYear = 2008 AND b.language = "English"
RETURN b.bookId
```

Q5. Provide all the review information for the book with ID 1598428, including the users who wrote them.

```
MATCH (:Book{bookId:1598428})-[:Has]->(r:Review)<-[:Writes]-(u:User)
RETURN r, u.userId
```

Q6. Given the user ID 24259, determine all the information about the books they read and their genres.

```
MATCH (:User{userId:24259})-[:Reads]->(b:Book)-
      [:categorized_As]->(g:Genre)
RETURN b, g
```

Q7. Provide all the books published in the last year along with their publishers.

```
MATCH (b:Book{publicationYear:2024})<-[:Publishes]-(p:Publisher)
RETURN b.bookId, p.publisherId
```

Q8. Show 10 books for the home page.

```
MATCH (b:Book)
RETURN b.bookId
LIMIT 10
```

8 Why Cassandra?

According to what we have already said from the beginning of this document, most of all considering point 2 and point 4, we are interested in availability more than consistency (an eventual consistency level is enough) and we decided to use Cassandra as our database for our application.

Cassandra, according to the CAP theorem, is focused on partition tolerance and availability, provided by its P2P structure, both key features of our application. It is the best suite for a read/write intensive application such as our library platform, and with its consistent hashing, it offers us a very high scalability to scale our platform for more and more new users. Talking about the workload, it might be quite complex for the Cassandra system because of its limits in querying, however, thanks to the speed in resolving queries, its high write throughput, fault tolerance, high availability, and the possibility of filtering the result at application level, Cassandra is the best choice for our application.

MongoDB is more flexible in solving queries because we can use nested attributes and lookups and there are fewer limitations concerning Cassandra. In fact, at the application level, everything is visible in Mongo.

Despite that, MongoDB is focused on consistency and partition tolerance (CP) and this is not what our application needs.

Neo4j is very fast and flexible and allows us to make complex queries with Cypher which is declarative and provides clauses to filter, return, create, delete, and set data; it is very fast thanks to its traversal speed, due to its index-free adjacency, where relationships between nodes are persistent and not calculated at query time.

Despite that, since Neo4j is well suited mostly for read-intensive and analytical applications, it is not the best database structure for our application, in fact, we need also to manage lots of writes from users.

For sure, it will be the best fit for our possible future implementation as we have already said in point 1: a recommendation system for our platform.

9 Cassandra configuration

For deploying our database, we relied on the docker features: we created a 3-node cluster using docker-compose.

The Docker Compose file sets up a Cassandra cluster with three nodes for fault-tolerant distributed database management with some key features:

- Persistent storage is defined for each node (cassandra-node-1, cassandra-node-2, cassandra-node-3) to ensure data durability even if containers are restarted.
- Each node has a health check that uses `nodetool status` to verify that the node is operational.
- The health check ensures dependent services (e.g., cassandra-2 depends on a healthy cassandra-1) only start when the required services are ready.
- The `restart: on-failure` directive ensures that any node will automatically restart if it crashes, contributing to fault tolerance by minimizing downtime.

The replication strategy is defined at the keyspace level. During the creation of keyspace, we can define the class and the replication factor that we want to use. In our case, we have decided to use the `SimpleStrategy` with a replication factor equal to 3 (the data are replicated in all three nodes) to ensure the maximum availability level.

10 Cassandra logical schema

10.1 Keyspace

```
CREATE KEYSPACE IF NOT EXISTS library
WITH REPLICATION = {'class':'SimpleStrategy', 'replication_factor':'1'};
```

10.2 Types

Three different types are used to perform queries efficiently:

- Review_t: [6.1.1 Book aggregate](#)
- Book_at: [6.1.2 Author aggregate](#)
- Book_ut: [6.1.3 User aggregate](#)

10.3 Tables

Four different tables to store efficiently all the data:

- Book1: [6.1.1 Book aggregate](#)
- Book2: [6.1.1 Book aggregate](#)
- Author: [6.1.2 Author aggregate](#)
- User: [6.1.3 User aggregate](#)

10.4 Copy the dataset from CSV files

```
COPY library.book1 (bookId, title, publicationYear, description,
    language, reviews, publisherId)
FROM '/csv/books.csv'
WITH HEADER = TRUE;

COPY library.book2 (bookId, title, publicationYear, description,
    language, reviews, publisherId)
FROM '/csv/books.csv'
WITH HEADER = TRUE;

COPY library.author (authorid, books)
FROM '/csv/authors.csv'
WITH HEADER = TRUE AND MINBATCHSIZE=1 AND MAXBATCHSIZE=1
    AND PAGESIZE=10;

COPY library.user (userid, books)
FROM '/csv/users.csv'
WITH HEADER = TRUE AND MINBATCHSIZE=1 AND MAXBATCHSIZE=1
    AND PAGESIZE=10;
```

11 Cassandra instance

```
cqlsh:library> describe library

CREATE KEYSPACE library WITH replication = {'class': 'SimpleStrategy', 'replication_factor': '3'} AND durable_writes = true;

CREATE TYPE library.book_at (
  bookid int,
  publicationyear int,
  genres set<text>
);

CREATE TYPE library.book_ut (
  bookid int,
  title text,
  publicationyear int,
  description text,
  language text,
  genres set<text>
);

CREATE TYPE library.review_t (
  reviewid int,
  score int,
  comment text,
  date text,
  userid int
);

CREATE TABLE library.author (
  authorid int PRIMARY KEY,
  books set<frozen<book_at>>
) WITH additional_write_policy = '99p'
  AND allow_auto_snapshot = true
  AND bloom_filter_fp_chance = 0.01
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
  AND cdc = false
  AND comment = ''
  AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
  AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
  AND memtable = 'default'
  AND crc_check_chance = 1.0
  AND default_time_to_live = 0
  AND extensions = {}
  AND gc_grace_seconds = 864000
  AND incremental_backups = true
  AND max_index_interval = 2048
  AND memtable_flush_period_in_ms = 0
  AND min_index_interval = 128
  AND read_repair = 'BLOCKING'
  AND speculative_retry = '99p';

CREATE TABLE library.book1 (
  bookid int PRIMARY KEY,
  description text,
  language text,
  publicationyear int,
  publisherid int,
  title text,
  reviews set<frozen<review_t>>
) WITH additional_write_policy = '99p'
  AND allow_auto_snapshot = true
  AND bloom_filter_fp_chance = 0.01
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
  AND cdc = false
  AND comment = ''
  AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
  AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
  AND memtable = 'default'
  AND crc_check_chance = 1.0
  AND default_time_to_live = 0
  AND extensions = {}
  AND gc_grace_seconds = 864000
  AND incremental_backups = true
  AND max_index_interval = 2048
  AND memtable_flush_period_in_ms = 0
  AND min_index_interval = 128
  AND read_repair = 'BLOCKING'
  AND speculative_retry = '99p';

CREATE TABLE library.book2 (
  publicationyear int,
  language text,
  bookid int,
  description text,
  publisherid int,
  title text,
  reviews set<frozen<review_t>>,
  PRIMARY KEY (publicationyear, language, bookid)
) WITH CLUSTERING ORDER BY (language ASC, bookid ASC)
  AND additional_write_policy = '99p'
  AND allow_auto_snapshot = true
  AND bloom_filter_fp_chance = 0.01
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
  AND cdc = false
  AND comment = ''
  AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
  AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
  AND memtable = 'default'
  AND crc_check_chance = 1.0
  AND default_time_to_live = 0
  AND extensions = {}
  AND gc_grace_seconds = 864000
  AND incremental_backups = true
  AND max_index_interval = 2048
  AND memtable_flush_period_in_ms = 0
  AND min_index_interval = 128
  AND read_repair = 'BLOCKING'
  AND speculative_retry = '99p';

CREATE TABLE library.user (
  userid int PRIMARY KEY,
  books set<frozen<book_ut>>
) WITH additional_write_policy = '99p'
  AND allow_auto_snapshot = true
  AND bloom_filter_fp_chance = 0.01
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
  AND cdc = false
  AND comment = ''
  AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
  AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
  AND memtable = 'default'
  AND crc_check_chance = 1.0
  AND default_time_to_live = 0
  AND extensions = {}
  AND gc_grace_seconds = 864000
  AND incremental_backups = true
  AND max_index_interval = 2048
  AND memtable_flush_period_in_ms = 0
  AND min_index_interval = 128
  AND read_repair = 'BLOCKING'
  AND speculative_retry = '99p';
```

Figure 4: Library KEYSPACE

12 Workload implementation

12.1 Q1

```
sqlsh:library> SELECT * FROM Book1 WHERE bookid = 238639;
@ Row 1
+-----+
| bookid |
+-----+
| 238639 |
+-----+
| description |
+-----+
| Corrie ten Boon was a Dutch watchmaker who became a heroine of the Resistance a survivor of Hitler's concentration camps and one of the most remarkable evangelists of the twentieth century. In World War II she and her family risked their lives to help Jews and underground workers escape from the Nazis and for their work they were tested in the infamous Westerbork death camp. Only Corrie among her family survived to tell the story of how faith ultimately triumphs over evil. Here is the riveting account of how Corrie and her family were able to save many of God's chosen people. For 35 years millions have seen that there is no pit so deep that God's love is not deeper still. Now The Hiding Place repackaged for a new generation of readers continues to declare that God's love will overcome, heal and restore.
+-----+
| language |
+-----+
| English |
+-----+
| publicationyear |
+-----+
| 1971 |
+-----+
| publisherid |
+-----+
| null |
+-----+
| reviews |
+-----+
| ((reviewid: 371984, score: 5, comment: 'None', date: '2020-10-19', userid: 1672), (reviewid: 1307460, score: 5, comment: 'I think this book should be read by all school kids in addition to or in place of The Diary of Anne Frank', date: '2007-09-25', userid: 5688), (reviewid: 2080643, score: 5, comment: 'None', date: '2011-04-07', userid: 8746), (reviewid: 2448208, score: 3, comment: 'None', date: '2021-08-08', userid: 18200), (reviewid: 4335742, score: 5, comment: 'None', date: '2023-01-24', userid: 18884), (reviewid: 4824159, score: 5, comment: 'This is a great book about finding faith in truly disheartening times. Corrie ten Boon was a brave Dutch woman who along with her family worked with the Dutch underground to hide Jews from the German soldiers. They obtained stolen ration cards, forged papers and addresses for families in outlying areas who could take them in. The few that were a higher risk because they looked too Jewish or they had health problems stayed with Corrie and her family. A secret room was added to the house so they could hide when the Gestapo came. They had drills and tried to be prepared for the grueling questioning by the German soldiers but how can one really prepare for such a thing?', date: '2011-09-16', userid: 21337), (reviewid: 5134207, score: 5, comment: 'None', date: '2016-06-22', userid: 22207), (reviewid: 5762494, score: 3, comment: 'None', date: '2023-01-01', userid: 23948))
+-----+
| title |
+-----+
| The Hiding Place |
+-----+
(1 rows)
```

Figure 5: First query

12.2 Q2

```
sqlsh:library> SELECT books FROM Author WHERE authorId = 737939;
@ Row 1
+-----+
| books |
+-----+
| ((bookid: 1316372, publicationyear: 2014, genres: {}), (bookid: 1443567, publicationyear: 2014, genres: {'Adventure', 'Dystopian', 'Fantasy', 'Science fiction', 'Time Travel', 'War'}), (bookid: 1450222, publicationyear: 1992, genres: {'Cardinals', 'Classics', 'Fantasy'}), (bookid: 1450255, publicationyear: 1991, genres: {'Classics', 'Fantasy'}), (bookid: 1450268, publicationyear: 1990, genres: {'Adventure', 'Classics', 'Fantasy', 'War'}), (bookid: 1585290, publicationyear: 2012, genres: {'War'}), (bookid: 1585419, publicationyear: 2012, genres: {}))
+-----+
(1 rows)
```

Figure 6: Second query

12.3 Q3

```
sqlsh:library> SELECT books FROM Author WHERE authorId = 206880;
@ Row 1
+-----+
| books |
+-----+
| ((bookid: 187207, publicationyear: 2005, genres: {'Fantasy', 'Fiction', 'Science fiction', 'Suspense'}), (bookid: 544097, publicationyear: 2011, genres: {}), (bookid: 570600, publicationyear: 2003, genres: {}), (bookid: 570845, publicationyear: 2007, genres: {}), (bookid: 570851, publicationyear: 2003, genres: {}), (bookid: 570852, publicationyear: 2000, genres: {}), (bookid: 570862, publicationyear: 2010, genres: {}), (bookid: 570871, publicationyear: 2003, genres: {}), (bookid: 570935, publicationyear: 2011, genres: {}), (bookid: 716100, publicationyear: 1999, genres: {}), (bookid: 759938, publicationyear: 2008, genres: {}), (bookid: 808539, publicationyear: 2014, genres: {}), (bookid: 1212733, publicationyear: 2006, genres: {}))
+-----+
(1 rows)
```

Figure 7: Third query

12.4 Q4

```
cqlsh:library> SELECT bookId FROM Book2 WHERE publicationYear = 2017 AND language = 'English';
```

@ Row 1	
bookid	108266
@ Row 2	
bookid	440029
@ Row 3	
bookid	563864
@ Row 4	
bookid	601076
@ Row 5	
bookid	602456
@ Row 6	
bookid	676570
@ Row 7	
bookid	788492
@ Row 8	
bookid	798508
@ Row 9	
bookid	1100015
@ Row 10	
bookid	1151162
@ Row 11	
bookid	1151216
@ Row 12	
bookid	1151311
@ Row 13	

Figure 8: Fourth query

12.5 Q5

```

@ Row 1

review | [[reviewId: 1324549, score: 4, comment: 'None', date: '2011-04-09', userId: 5685], [reviewId: 1352759, score: 4, comment: 'p!ana Wynne Jones has been one of my favourite authors since I was a child but this book really highlighted just how much wit she had and how much of it she probably had to hide since she most well know for her childrens books This book is a parody of the entire Fantasy genre prebts this was published in 1996 with the tone of a very disapproving author br ltb!ocknotqetg!FELLOW TRAVELLERS These are people who join the Tour for a short while and then leave or get killed If they have NAMES and characters then you will be sorry to lose them otherwise not!tb!ocknotqetg!br This book made me laugh out loud with delight but also frequent ly tempted me to skin A huge negative of the book is its structure that of an A-Z glossary of terms commonly used in Fantasy and D&Gs hilarious definitions of them It seems to act as a sort of a brochure for you the reader who is also the Tourist of Fantasyland and introduce you to the sights and sounds and people youll meet on your Tour!br ltb!ocknotqetg!GAY Magic may be one of your COMPANIONS on the Tour he will be very beautiful and he will dress in gorgeous colours!r!tb!ocknotqetg!br While this is a pretty fun premise it got pretty tiring to read disconnected A-Z terms, especially when a lot of them refer to another term from a different letter usually denoted with all caps It could also be a function of me reading this on an e-reader and therefore unable to flip around like I could with a physical book and perhaps that had been really how DMO had intended for it to be read As it is I though my experience with it just felt very disjointed For example!br ltb!ocknotqetg!Baths are the occasion for SEX with one or more of your FELLOW TRAVELLERS No matter how irritating you have found her!ia up to then after or during the Bath you will find her!ia irremissible It is probably something in the MATERIAL!tb!ocknotqetg!br I kept thinking about how this book would be so much more fun if it had been sectioned off by themes and categories Instead of being by alphabets for example a whole section of Worlds or Settings and then another on Peoples and Cultures It could still follow the same premise of being a sort of travel guidebook for Tourists of Fantasyland but having a more coherent line of thought as we go along!br However I thoroughly appreciated how much DMO called out the misogyny that was pretty rampant and typical of the fantasy genre of the time!br ltb!ocknotqetg!It is clear that only males get to visit the human islands and only females find their way to the Oregon This is the tough sexist way the Manag ement wrote the Rules!r!tb!ocknotqetg!br ltb!ocknotqetg!DARK LADY There is never one of these so see DARK LADY instead!r!tb!ocknotqetg!br This book was also clearly heavily influenced by the Tolkien brand of high fantasy and I felt like a lot of entries were specifically poking fun at Lord of the Rings In sure DMO did this in a respectful and affectionate way though seeing that she st udied under Tolkien at Oxford when he was in the midst of writing Lord of the Rings I also definitely felt like I could see her tropes applying to other series like Wheel of Time and even the Harry Potter Trilogy though that is much more recent!br ltb!ocknotqetg!GIRLS or KINGS and SHORTS You have problems Rings have to be returned whence they came preferably at over a thousand degrees Celcius in the Core means you wont want to do this!r!tb!ocknotqetg!br Nevertheless a great book for any fantasy lovers to check out especially if youre more familiar with traditional fantasy tropes and conventions from pre1990s titles Get a physical copy if you're able to and you might want to save this as a side book to enjoy once in a while!p, date: '2020-11-26', user!d: 5823}, [reviewId: 1361834, score: 2, comment: 'shell it should be 2 and 12 stars but as usual I rounded the rating up Hes some good bits like the entry about horses but gets a bit repetitive after a while p, date: '2011-03-27', user!d: 8078}, [reviewId: 1899147, score: 4, comment: 'None', date: '2012-02-20', user!d: 6888}, [reviewId: 2042426, score: 4, comment: 'None', date: '2008-06-19', user!d: 8894}, [reviewId: 3270162, score: 4, comment: 'None', date: '2013-01-18', user!d: 13825}, [reviewId: 4488898, score: 3, comment: 'None', date: '2015-12-30', user!d: 19573}, [reviewId: 5958896, score: 4, comment: 'None', date: '2025-01-20', user!d: 26960}}]

(1 rows)

```

Figure 9: Fifth query

12.6 Q6

```

cqlsh:library> SELECT books FROM User WHERE userId = 54;

@ Row 1
books | ((bookid: 130, title: 'Dreams from My Father A Story of Race and Inheritance', publicationyear: 1995, description: 'Before Barack Obama became a politician he was among other things a w
riter. Dreams from My Father is his masterpiece, a refreshing revealing portrait of a young man asking the big questions about identity and belonging. The son of a black African father and a white
American mother, Obama recounts an emotional odyssey. He retraces the migration of his mother's family from Kansas to Hawaii then to his childhood home in Indonesia. Finally he travels to Kenya wher
e he confronts the bitter truth of his father's life and at last reconciles his divided inheritance', language: 'English', genres: ('African Americans', 'Biography', 'Autobiography', '
History', 'Nonfiction', 'social science')), (bookid: 233, title: 'The King Leopolds Ghost', publicationyear: 1998, description: 'With an introduction by Barbara Kingsolver a compelling account of th
e Congo Massacre a holocaust which resulted in the deaths of millions of peopleIn Hochschild's groundbreaking book he explores the devastation and exploitation of the Congo Free State by King Leo
pold II of Belgium a study replete with cruel leaders corruption and a few genuine heroes. King Leopolds Ghost was awarded the Duff Cooper Prize in 1999', language: 'English', genres: ('Anti-slave
ry movements', 'Classics', 'Congo (Democratic Republic)', 'General', 'History', 'Politics', 'War')), (bookid: 380, title: 'Watchmen', publicationyear: 1987, description: 'Considered the greatest
graphic novel in the history of the medium the Hugo Award-winning story chronicles the fall from grace of a group of superheroes plagued by altohuman failings Along the way the concept of the
superhero is dissected as an unknown assassin stalks the erstwhile heroes', language: 'English', genres: ('etc', 'Adventure', 'Classics', 'Comics', 'Comics & Graphic Novels', 'Dystopian', 'Fant
asy', 'Fiction', 'Graphic novels', 'Science fiction')), (bookid: 337, title: 'Becoming', publicationyear: 19, description: 'FM A LIFE filled with meaning and accomplishment Michelle Obama has an
argued as one of the most iconic and compelling women of our era As First Lady of the United States of Americathe first African American to serve in that role she helped create the most welcoming
and inclusive White House in history while also establishing herself as a powerful advocate for women and girls in the US and around the world dramatically changing the ways that families pursue
healthier and more active lives and standing with her husband as he led America through some of its most harrowing moments Along the way she showed us a few dance moves crushed Carpool Karaoke
and raised two downhearted daughters under an unforgiving media glareIn her memoir a work of deep reflection and mesmerizing storytelling Michelle Obama invites readers into her world chronickin
g the experiences that have shaped her from her childhood on the South Side of Chicago to her years as an executive balancing the demands of motherhood and work to her time spent at the worlds mo
st famous address with unerring honesty and lively wit she describes her triumphs and her disappointments both public and private telling her full story as she has lived it in her own words and o
n her own terms Warm wise and revelatory Becoming is the deeply personal reckoning of a woman of soul and substance who has steadily defied expectations and whose story inspires us to do the same
source: 1 https://becomingmichelleobama.com', language: 'English', genres: ('African American women lawyers', 'Biography', 'Biography & Autobiography', 'Juvenile Nonfiction', 'Nonfiction', 'Self-He
lp', 'Women', 'Young Adult', 'autobiography'))

[1 rows]

```

Figure 10: Sixth query

12.7 Q7

```

cqlsh:library> SELECT bookId, publisherId FROM book2 WHERE publicationYear = 2024;

@ Row 1
-----
bookid | 1245827
publisherid | null

@ Row 2
-----
bookid | 1245835
publisherid | null

@ Row 3
-----
bookid | 1245840
publisherid | null

@ Row 4
-----
bookid | 1289421
publisherid | null

@ Row 5
-----
bookid | 1353640
publisherid | null

@ Row 6
-----
bookid | 1409354
publisherid | 389

@ Row 7
-----
bookid | 1409356
publisherid | 107

@ Row 8
-----
bookid | 1464642
publisherid | null

@ Row 9
-----
bookid | 1556615
publisherid | 177342

@ Row 10
-----
bookid | 1565079
publisherid | 9592

@ Row 11

```

Figure 11: Seventh query

12.8 Q8

```
cqlsh:library> SELECT bookId FROM book1 LIMIT 10;
```

@ Row 1
bookid 978500
@ Row 2
bookid 1538144
@ Row 3
bookid 1464676
@ Row 4
bookid 1703353
@ Row 5
bookid 439535
@ Row 6
bookid 1151115
@ Row 7
bookid 1342760
@ Row 8
bookid 1556657
@ Row 9
bookid 798467
@ Row 10
bookid 1464681

(10 rows)

Figure 12: Eighth query

13 Conclusions

In this report, we explored the architecture of a distributed database system using Apache Cassandra. We defined a query model, designed and implemented dataset loading from CSV files, and discussed best practices for system implementation. Additionally, we examined using Docker to create a multi-node Cassandra cluster with replication for high availability and fault tolerance.

We also analyzed and implemented several queries for extracting data from the database, considering Cassandra's NoSQL characteristics. The solutions for handling complex queries, such as filtering and aggregation, were tailored to the system's specific features and business needs.

Future Perspectives:

As we already discussed, looking ahead, we could explore the implementation of a recommendation system using graph databases like Neo4j. This would allow for more efficient handling of relationships between users and items, enabling personalized recommendations based on complex graph structures.